

Kademlia – a P2P DHT

SSD 2022/23

FCUP-UP

Introduction

- Kademlia is a peer-to-peer (P2P) distributed hash table (DHT)
- Provides an overlay for building highly distributed applications
- Resilient to faults and attacks with no central points of failure (CPoF)
 - Including denial of service attacks
- Nodes can enter and leave the systems without compromising or disrupting the system

Kademlia Nodes

- Kademlia treats each node on a network as a leaf on a binary tree
 - Node have a 160-bit Unique Identifier (UId) (based on SHA-1)
 - A Node's position (in the tree) is determined by the shortest unique prefix of its UId
- Kademlia protocol ensures that each node knows of at least one node on each of its sub-trees
 - Thus, a node can locate any other node by its ID
- Node selection for storing a key-value pair relies on a notion of distance between two identifiers
 - Given two 160-bit identifiers, x and y , distance between them is $x \oplus y$ (XOR).

Inter-Node Messaging

- The Kademlia protocol consists of four Remote Procedure Calls (RPCs):
- **PING**: probes a node to see if it's online
- **STORE**: instructs a node to store a key-value pair
- **FIND_NODE**: returns information about the k nodes closest to the target id
- **FIND_VALUE**: similar to the FIND_NODE RPC, but if the recipient has received a STORE for the given key, it just returns the stored value

Node Joining

- Joining a Kademlia network requires discovery of a single node
 - Bootstrap node
- A joining node announces itself to the Bootstrap node
- The Bootstrap node replies with a set of “closest” nodes
 - Collects the NodeID from each response and adds it to its own peer table
- Parallel and asynchronous queries can prevent timeout delays or ‘retrieval hold-ups’
 - Due to failed nodes which have dropped off or left the network

Routing Tables and K-buckets


- The routing table is a binary tree whose leaves are **k-buckets**
 - Nodes maintain detailed knowledge of the address space closest to them
 - Exponentially decreasing knowledge of more distant address space
- **K-buckets** are lists of routing addresses of other nodes in the network
 - Contain the IP address, port, and NodeID of the respective nodes
- The routing table size is asymptotically bounded by $O(\log_2(n/k))$
 - Where n is the actual number of nodes in the network and k is the bucket size
 - Larger buckets slightly reduce the total number of buckets in the routing table.
- Symmetry is useful since it means that each of these closest contacts will be maintaining detailed knowledge of a similar part of the address space
 - Rather than a remote part.

Routing Tables and K-buckets

- Routing table is distributed and subject to different versions/visions
 - Each node maintains and manages a mapping for a subset of the nodes on the network in its own routing table
 - No absolute truth exists to where NodeIDs are mapped to their address


Join Network

- Each node has:
 - IP address
 - Port No
 - Unique ID
 - Tree<Buckets> (routing table)
 - Storage
- Special Nodes
 - Bootstrap Nodes
 - Know all nodes





Join Network

- Entering the group
 - Sends (IP, Port, ID) to a Bootstrap Node
 - Using FIND_NODE
 - Bootstrap Adds the new Node info to its routing table
 - Replies with a set of near nodes
 - Send (IP, Port, ID) to each nodes received
 - Repeat if new near nodes are received




Join Network

00110	
11010	

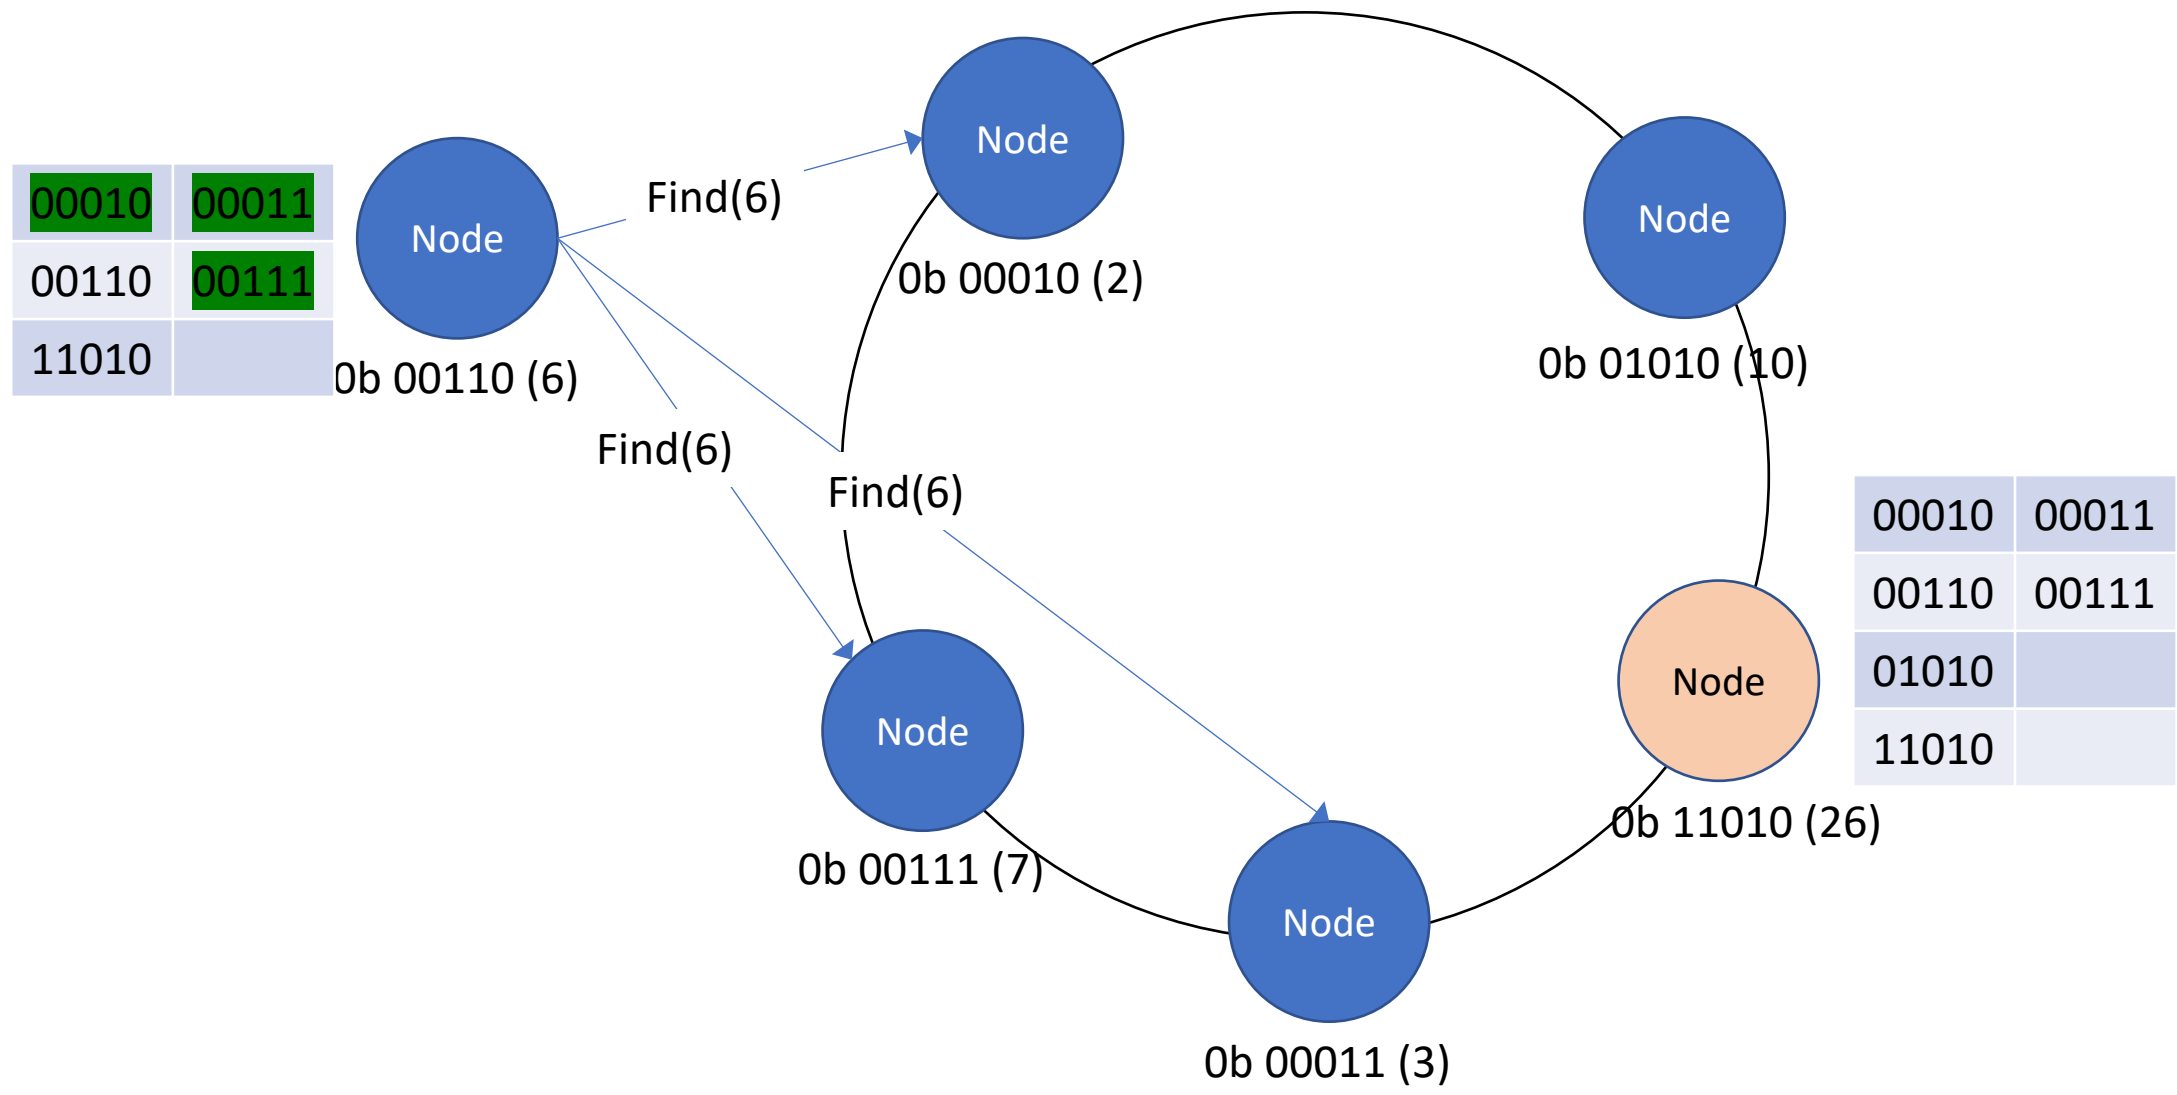


00010	00011
00111	
01010	
11010	

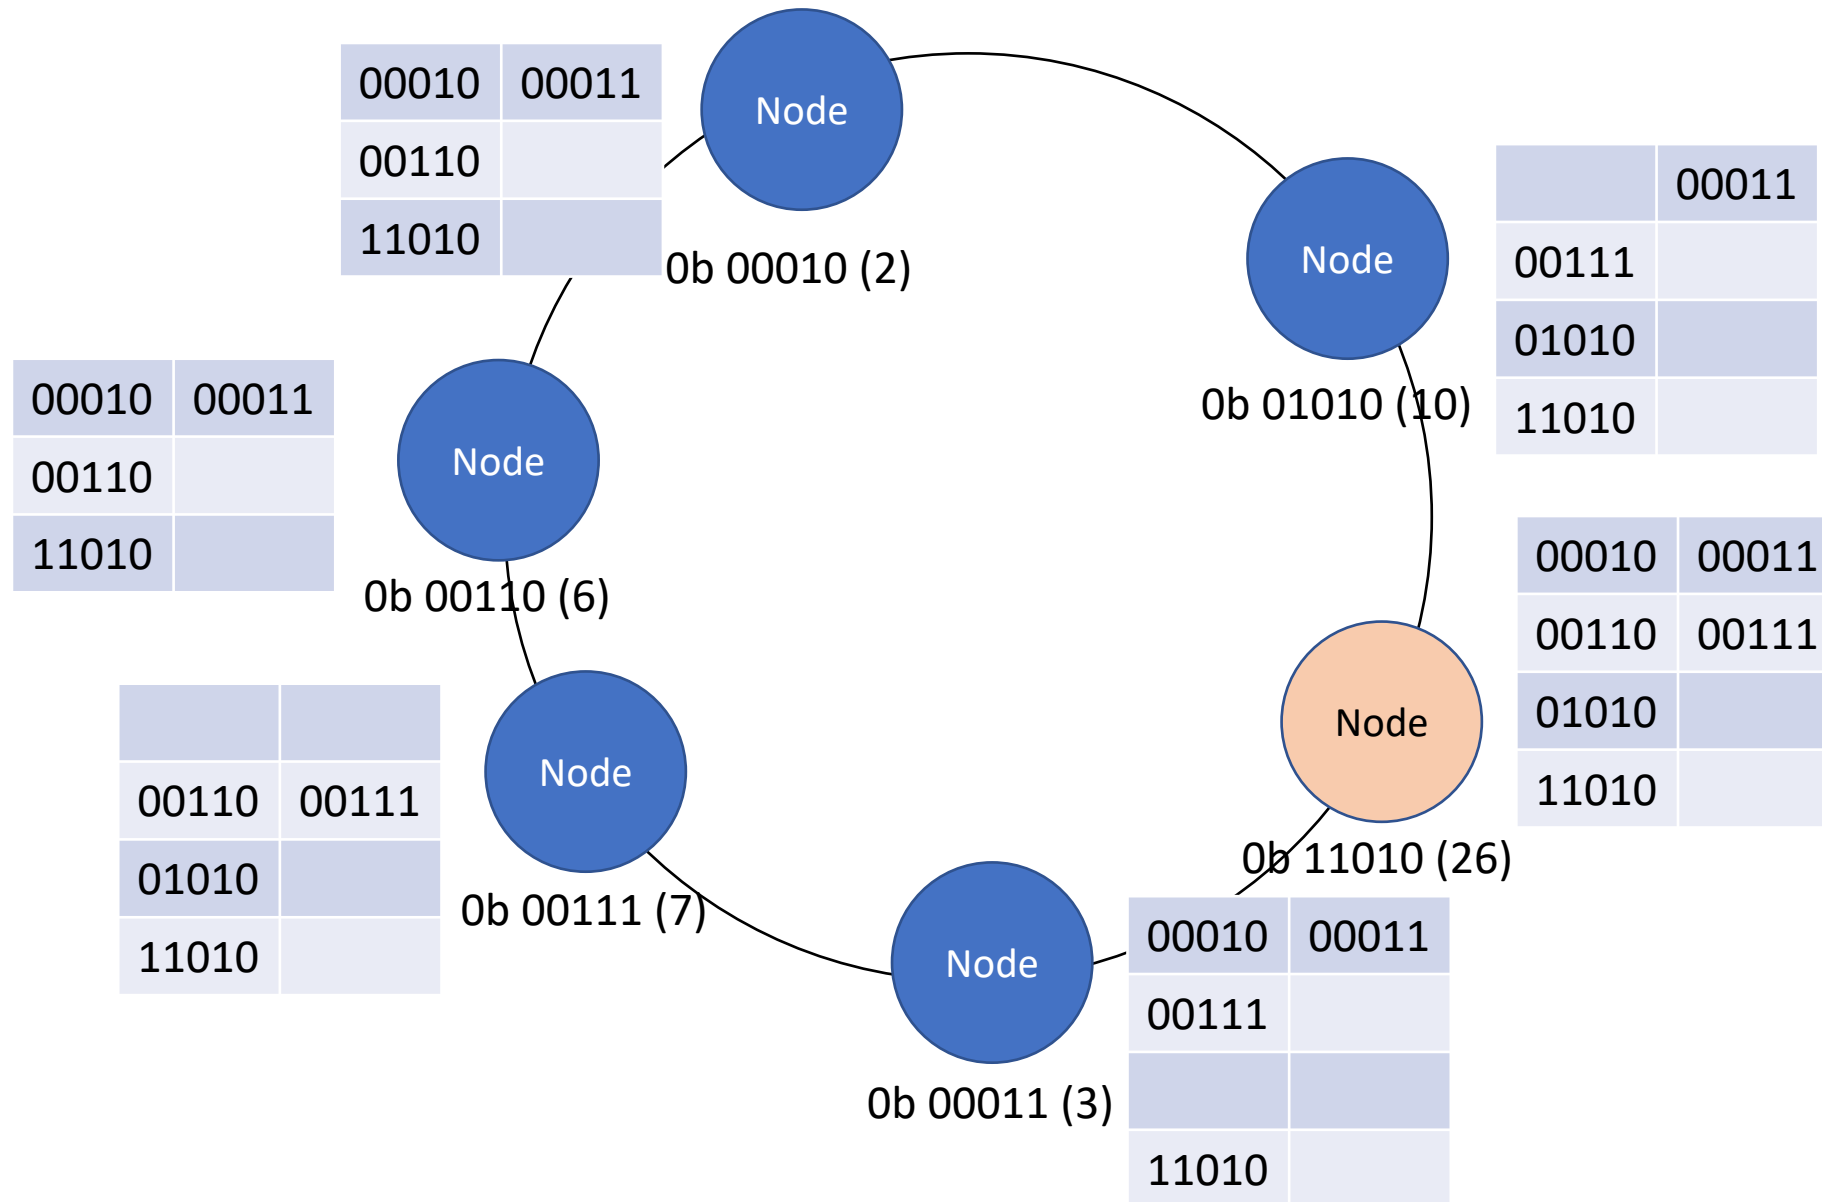
Join Network



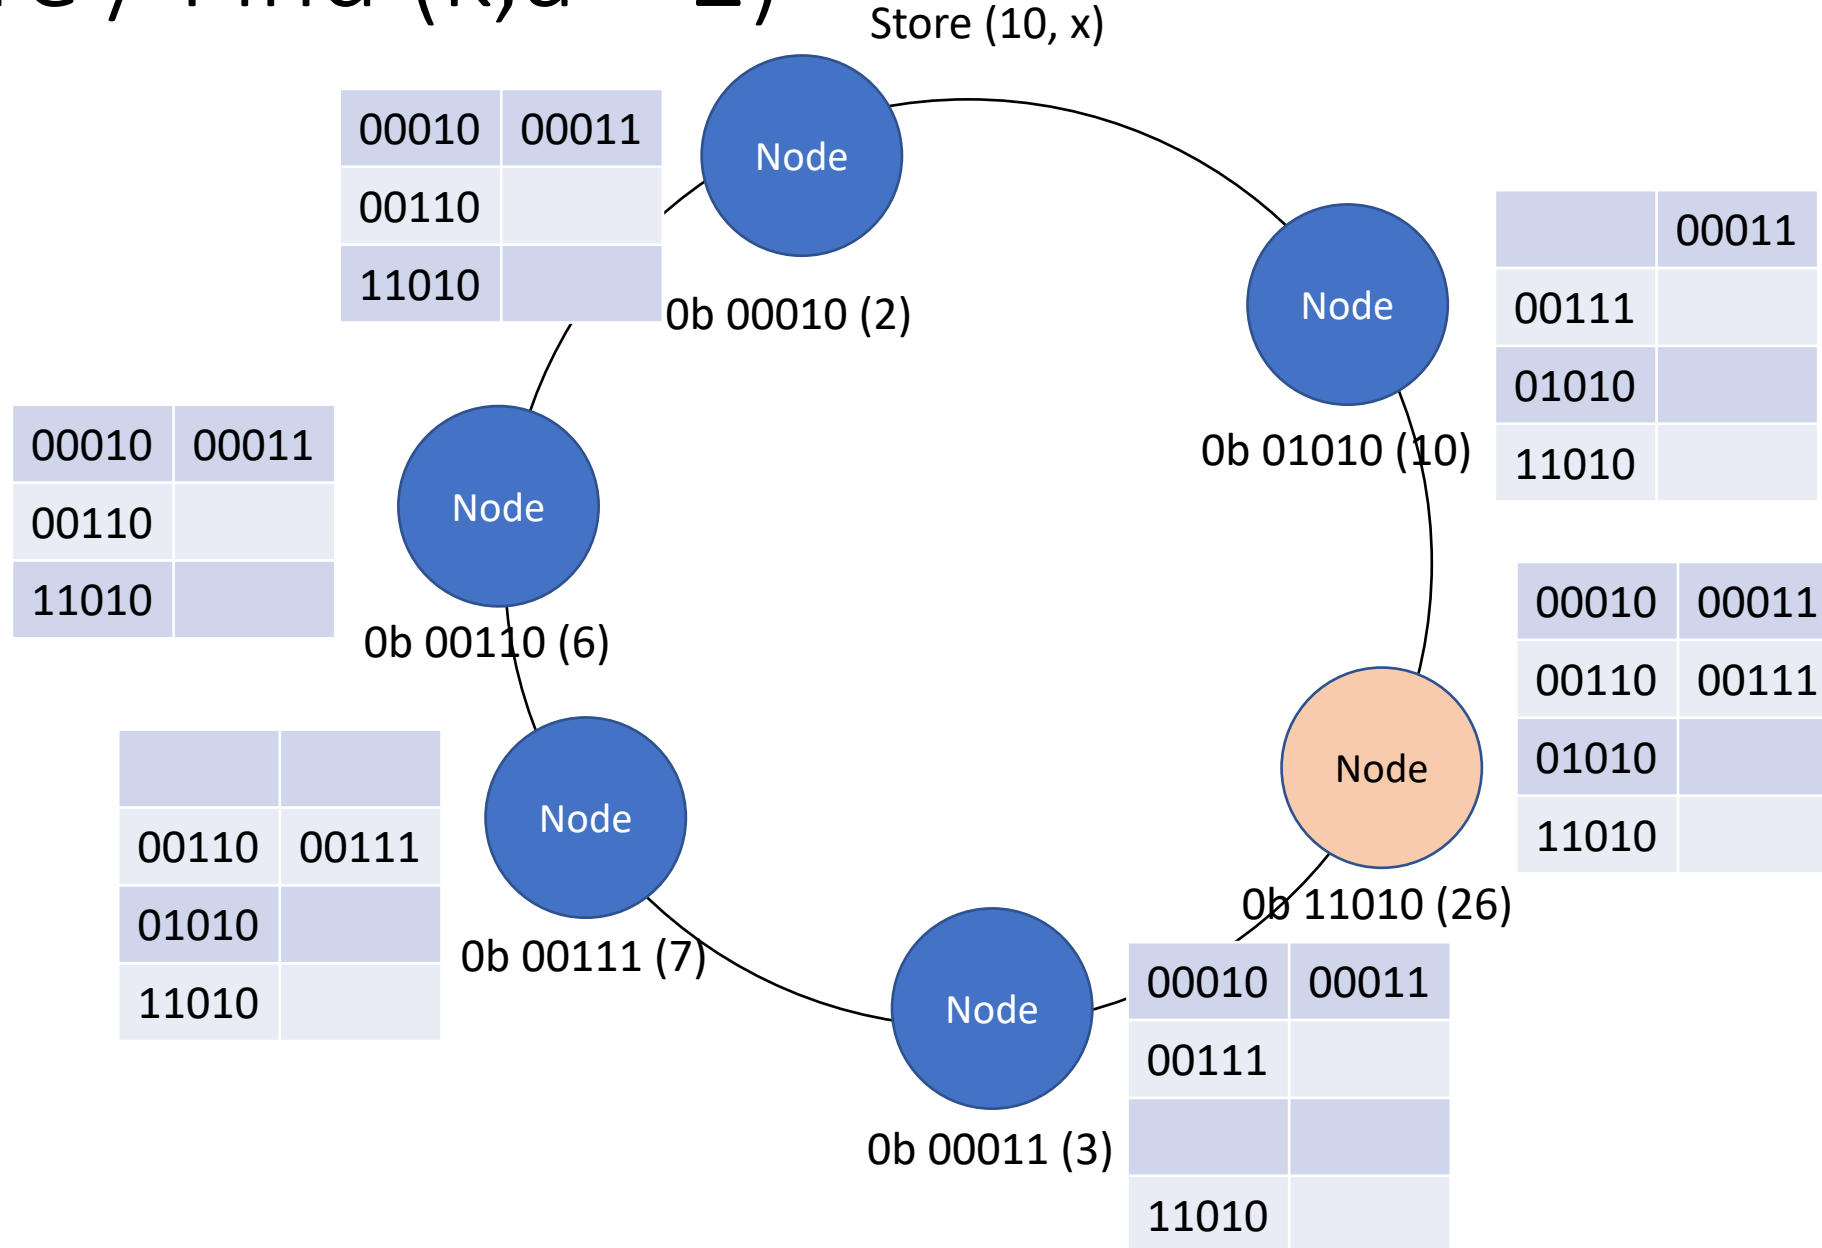
Join Network



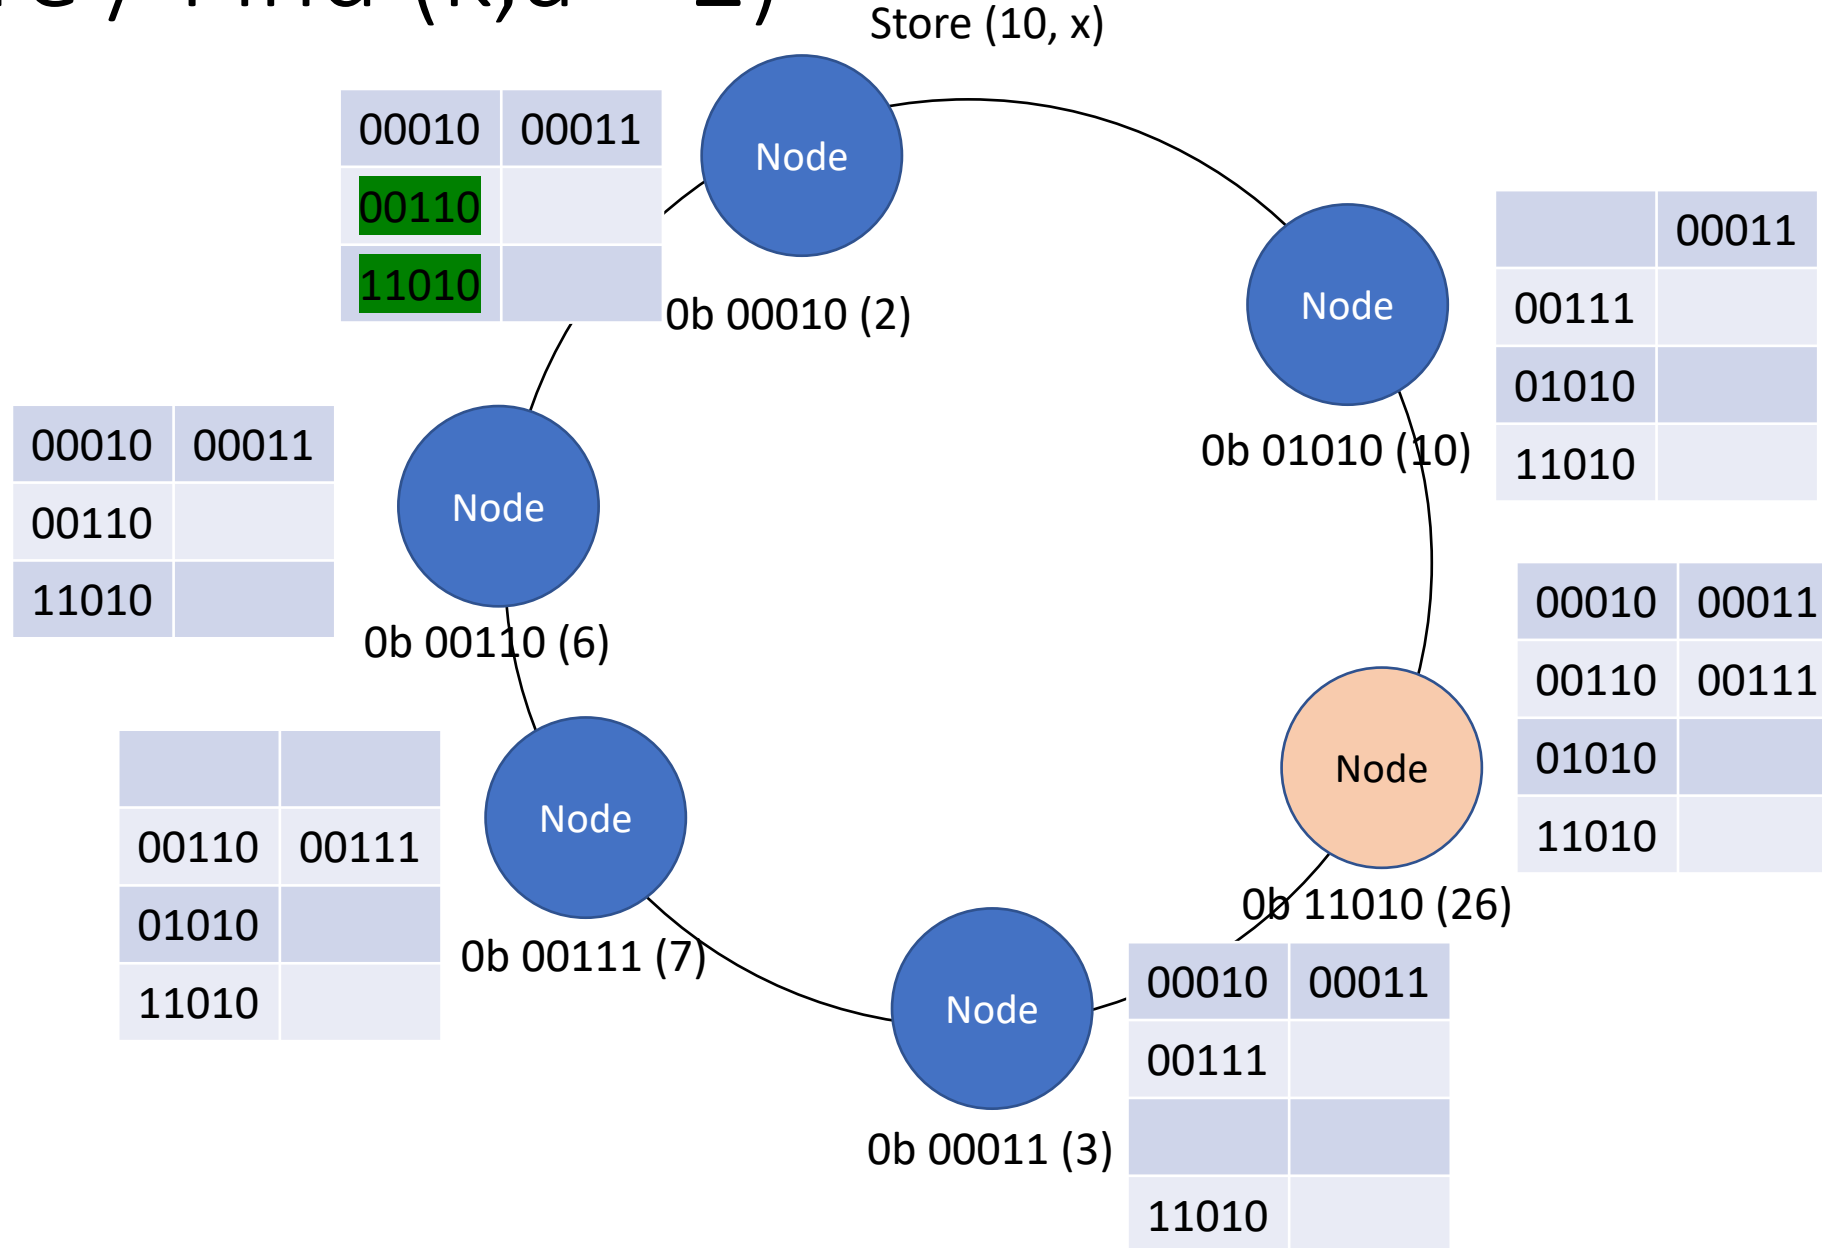
Store / Find



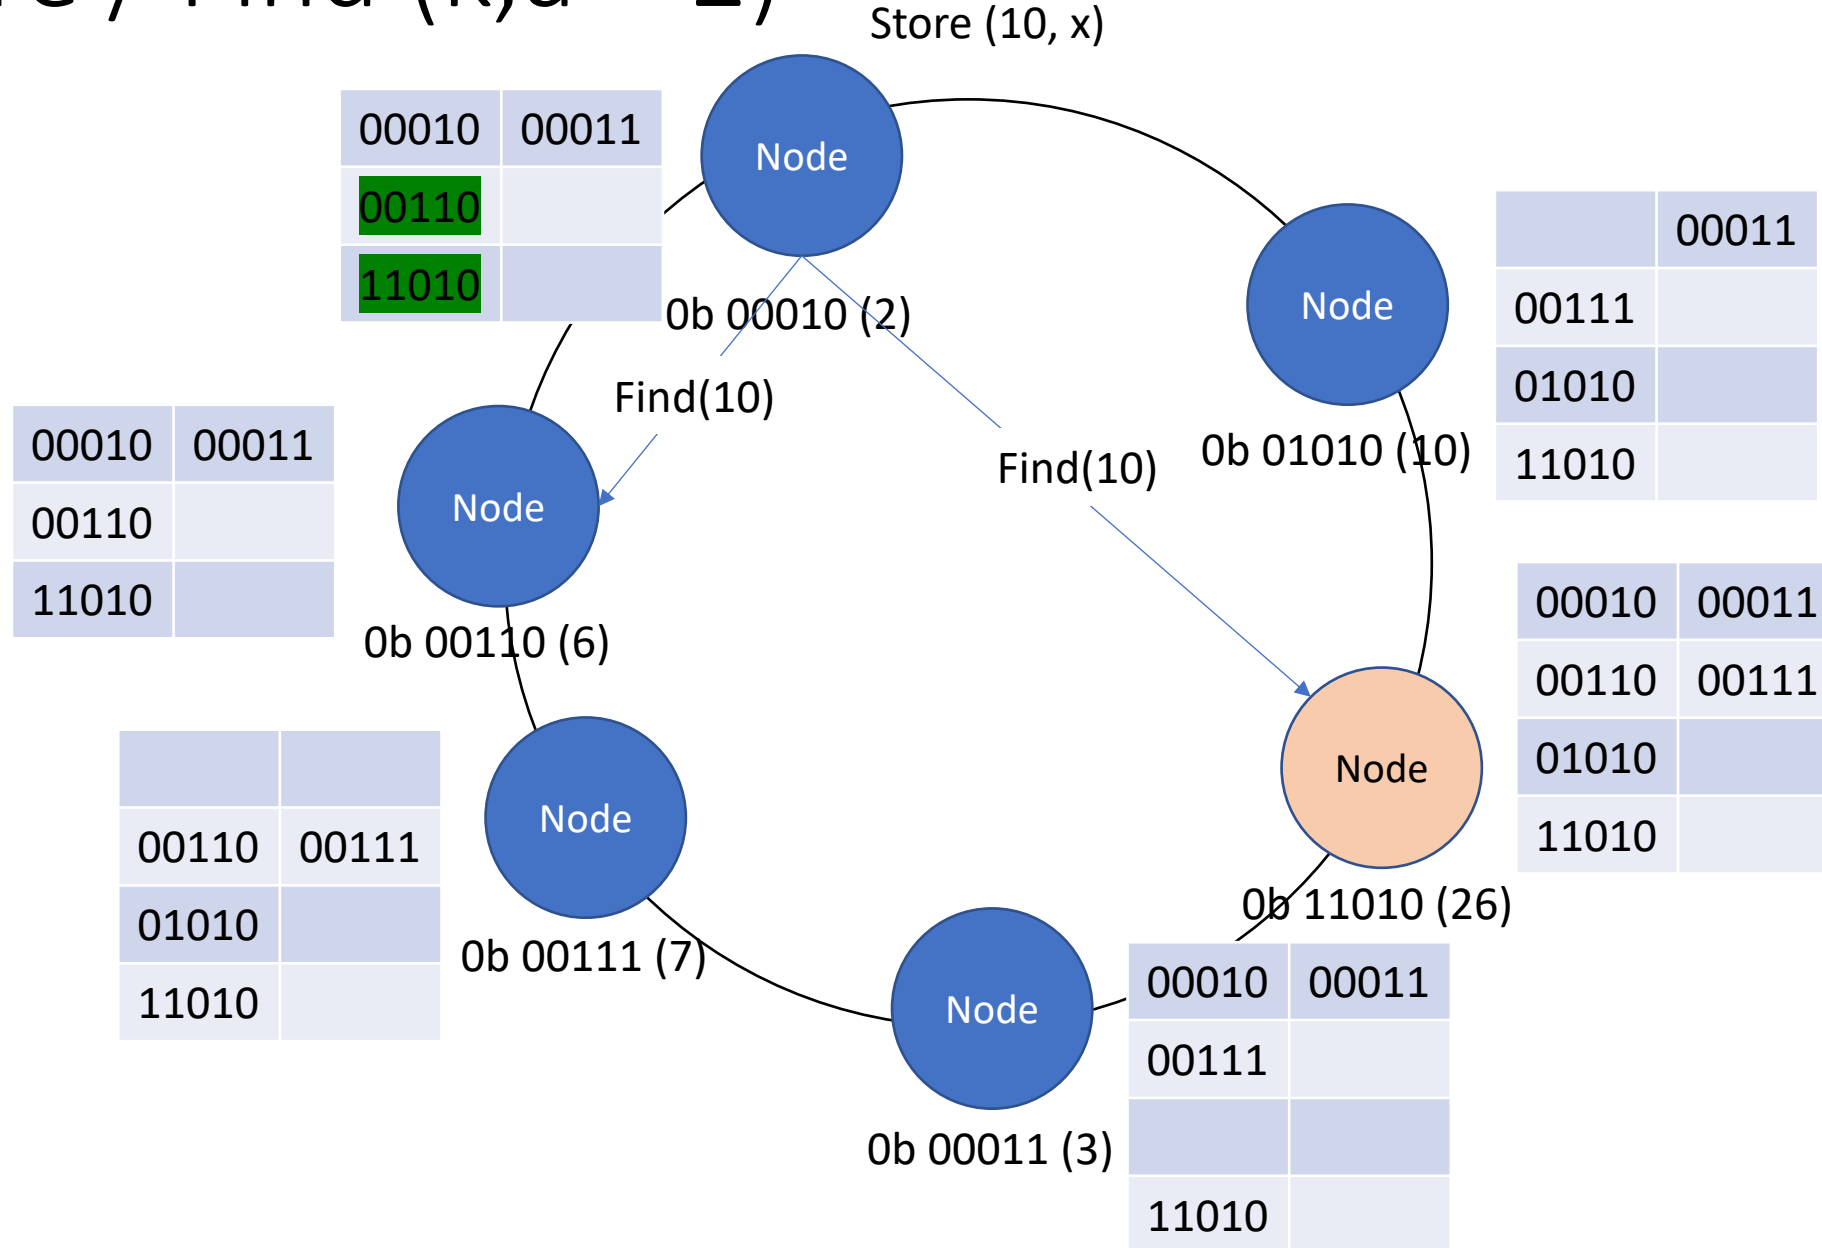
Store / Find (k,a = 2)



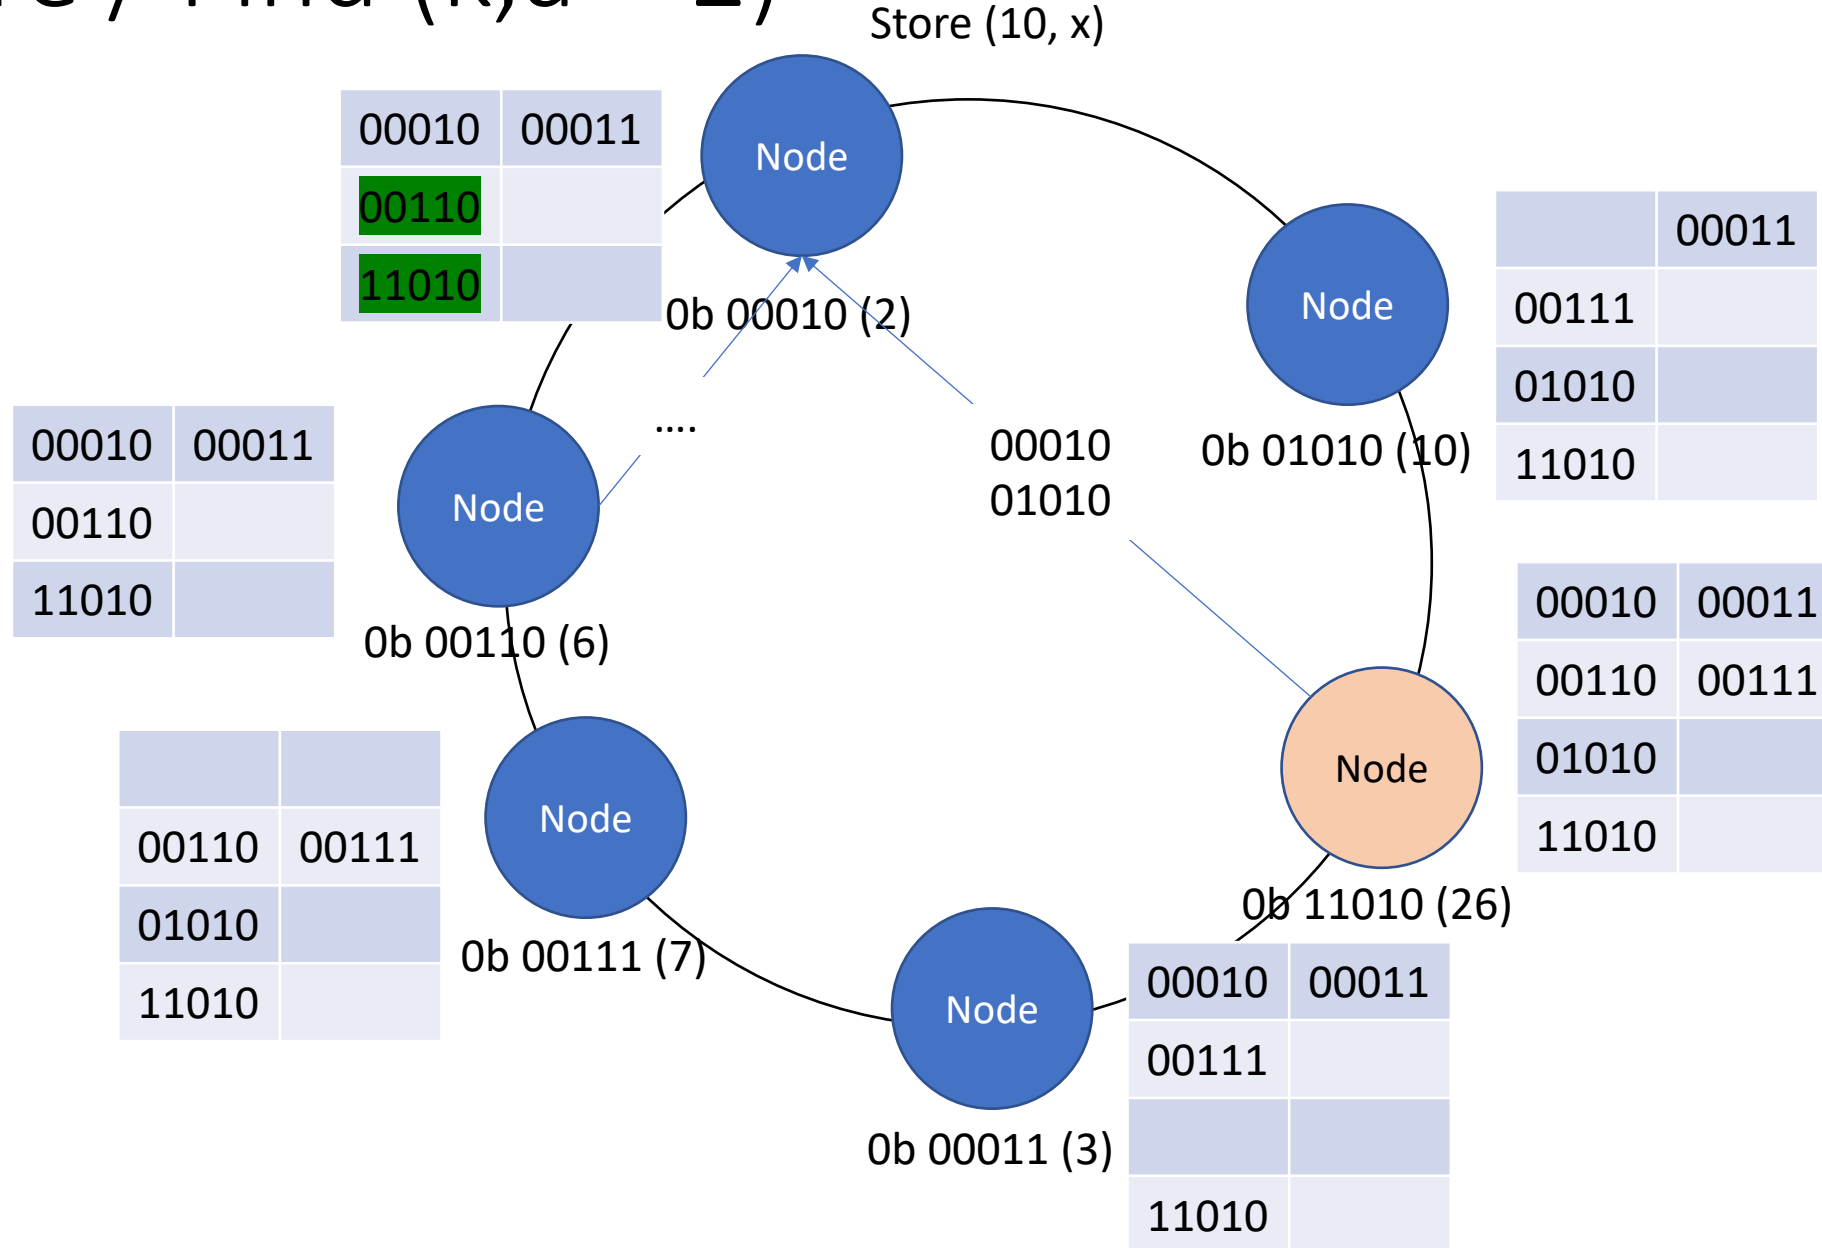
Store / Find (k,a = 2)



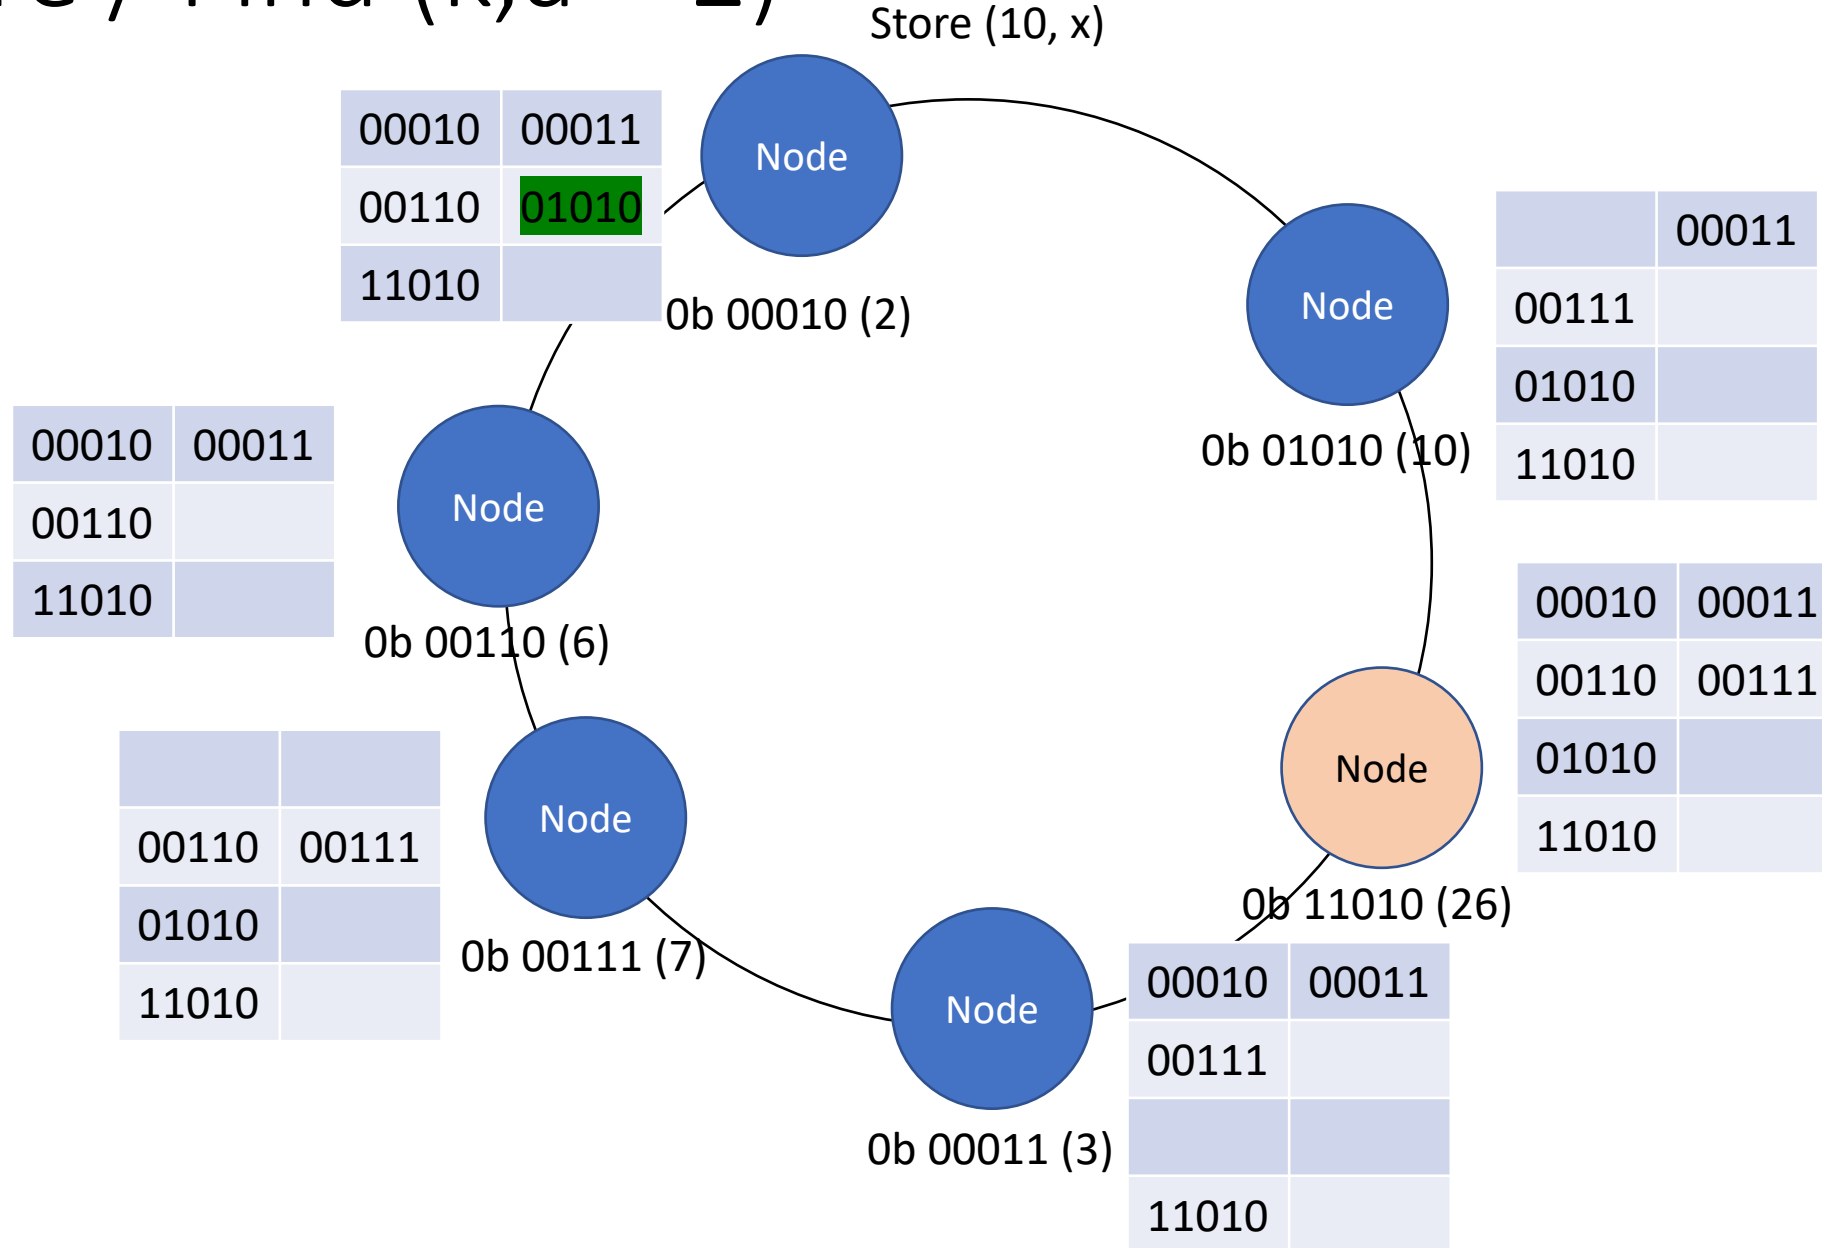
Store / Find (k,a = 2)



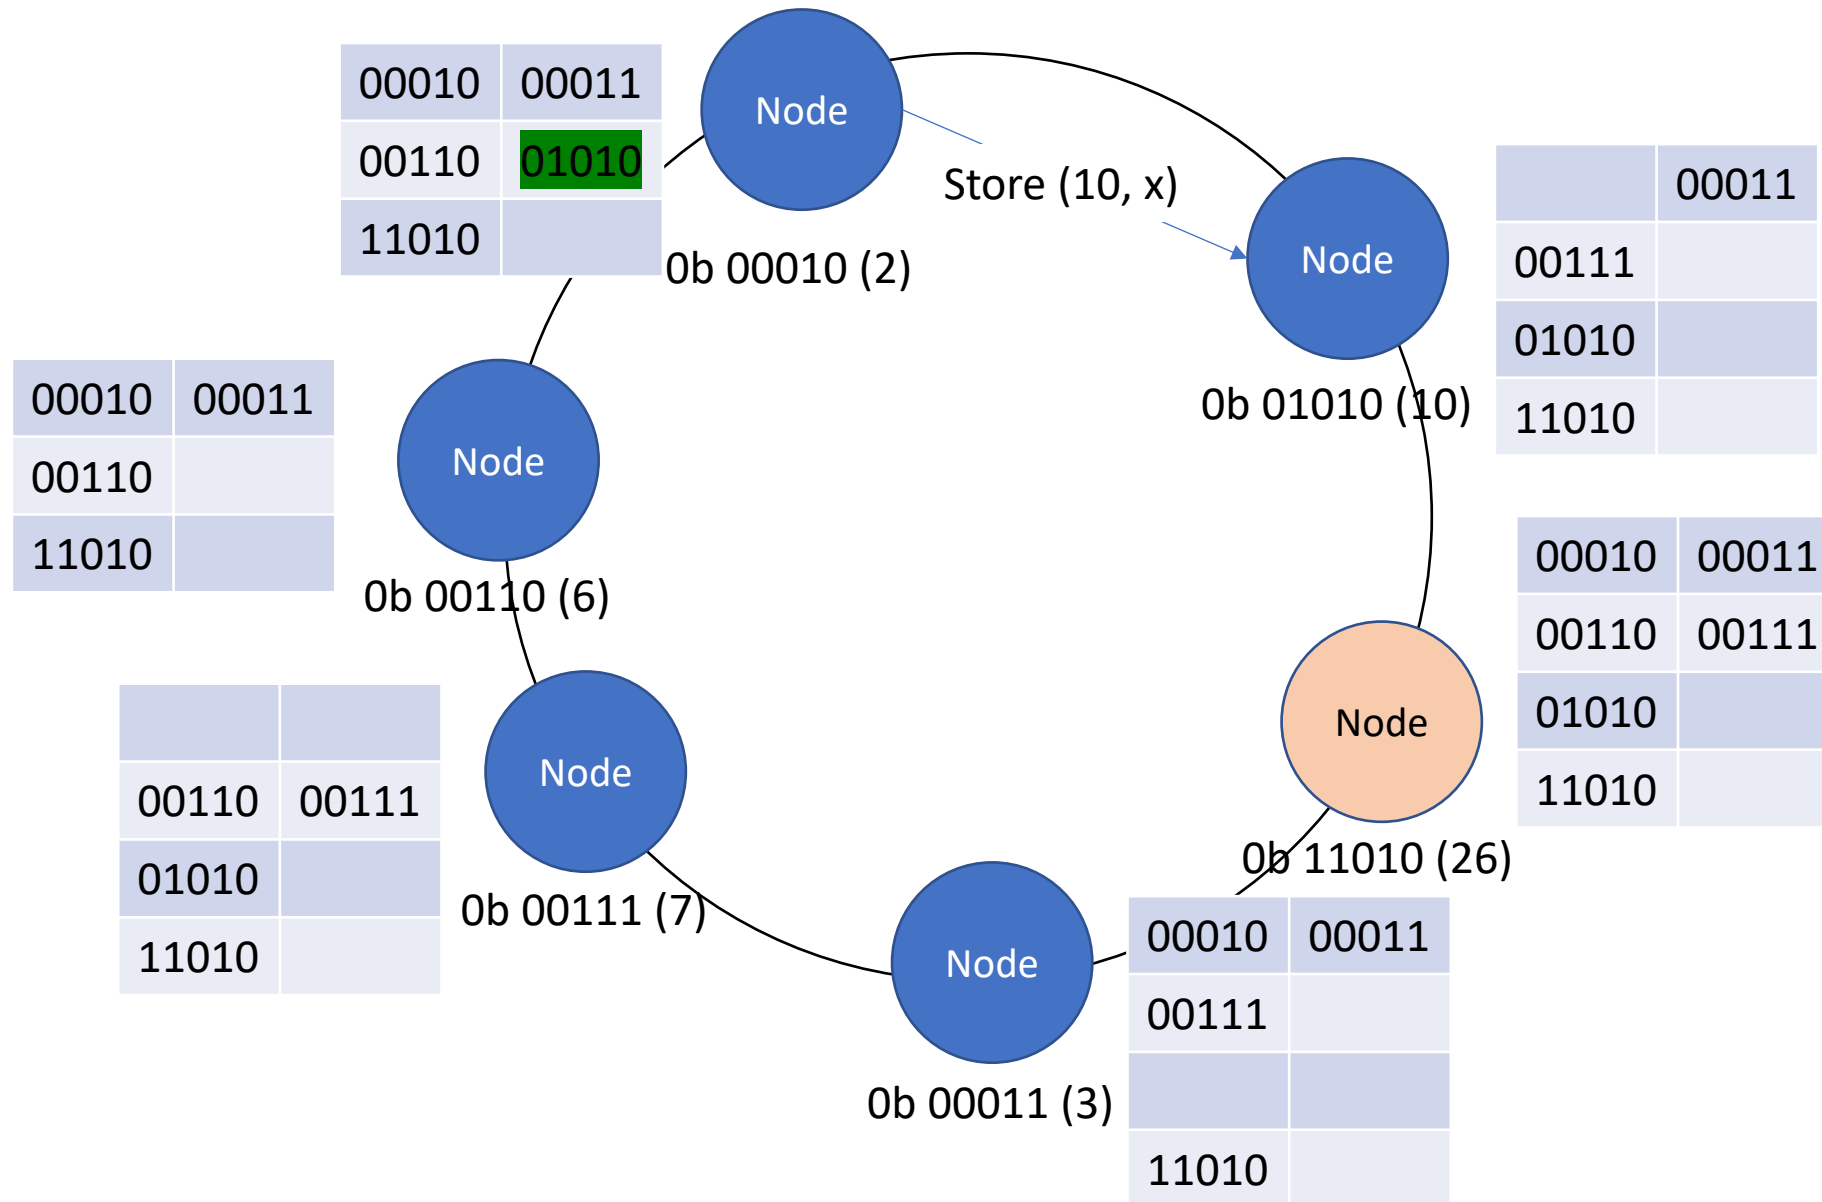
Store / Find (k,a = 2)



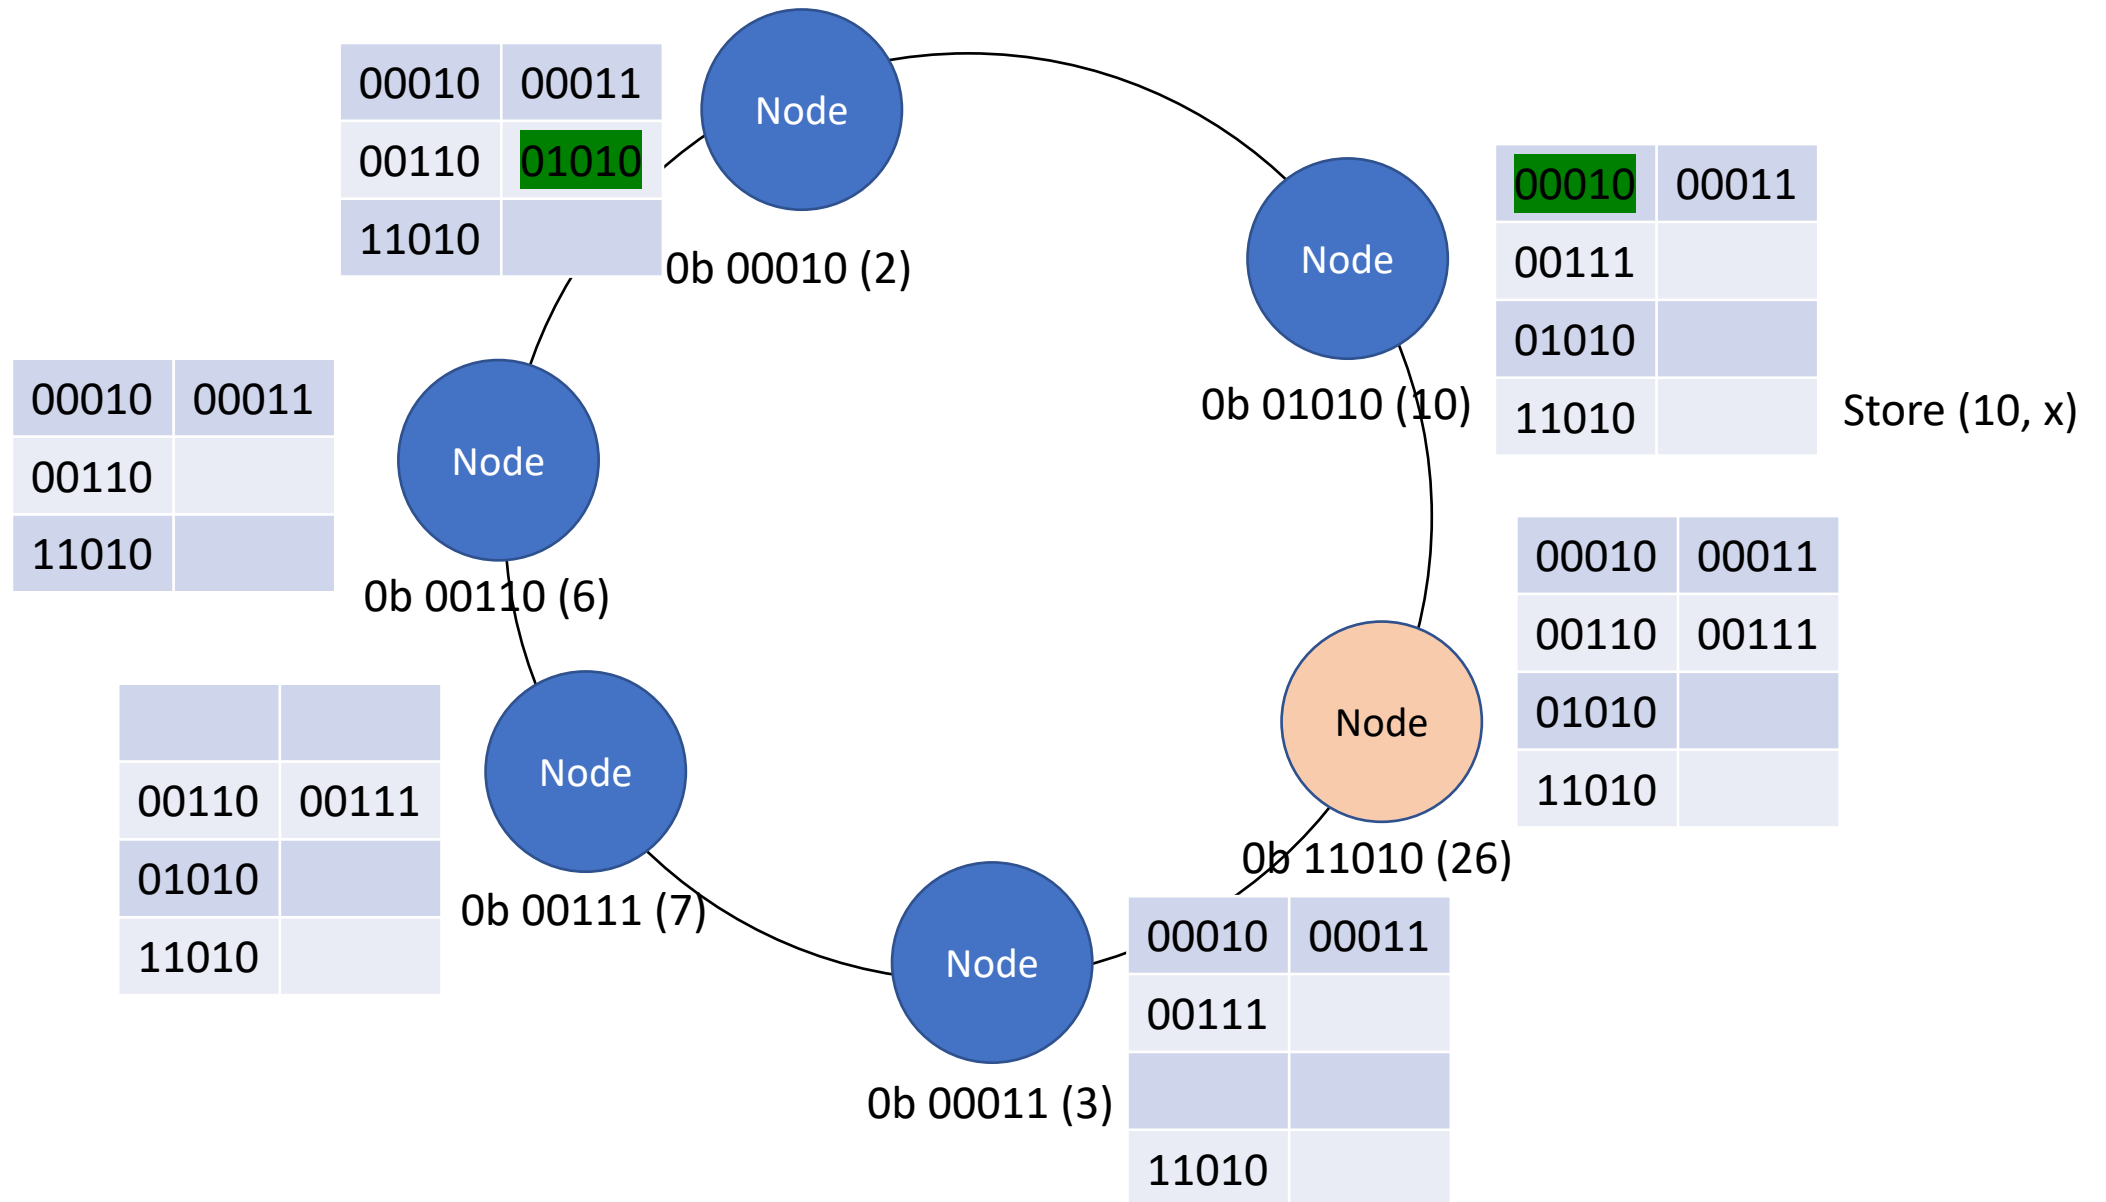
Store / Find (k,a = 2)



Store / Find (k,a = 2)



Store / Find (k,a = 2)



Store / Find (what ifs)

- What if a node fails to respond?
- What if a k-bucket is full?
- What if more than one node have the same distance?
- What if ... ?

References

- [0]: Maymounkov, P., & Mazières, D. (2002). Kademlia: A Peer-to-Peer Information System Based on the XOR Metric. In P. Druschel, F. Kaashoek, & A. Rowstron (Eds.), Peer-to-Peer Systems (pp. 53–65). Berlin, Heidelberg: Springer Berlin Heidelberg. https://doi.org/10.1007/3-540-45748-8_5
- [1]: Baumgart, I., & Mies, S. (2014). S / Kademlia : A Practicable Approach Towards Secure Key-Based Routing, (June). <https://doi.org/10.1109/ICPADS.2007.4447808>
- [2]: Freedman, M. J., & Mazières, D. (2003). Sloppy Hashing and Self-Organizing Clusters. In IPTPS. Springer Berlin / Heidelberg. Retrieved from www.coralcdn.org/docs/coral-iptps03.ps
- <https://github.com/libp2p/specs/tree/master/kad-dht>
- https://codethechange.stanford.edu/guides/guide_kademlia.html
- https://kelseyc18.github.io/kademlia_vis/basics/1/