

Postgres and Python for the visualization of Spatial Data

Advanced Topics in Databases

Psycopg – PostgreSQL database adapter for Python

- Psycopg is the most popular PostgreSQL database adapter for the Python programming language.
- Its main features are the complete implementation of the Python DB API 2.0 specification and the thread safety (several threads can share the same connection).
- It was designed for heavily multi-threaded applications that create and destroy lots of cursors and make a large number of concurrent INSERTs or UPDATEs.
- Psycopg 2 is mostly implemented in C as a libpq wrapper, resulting in being both efficient and secure.
- Many Python types are supported out-of-the-box and adapted to matching PostgreSQL data types;

The psycopg2 module content

- The module interface respects the standard defined in the DB API 2.0.

```
psycopg2.connect(dsn=None, connection_factory=None, cursor_factory=None,  
                async=False, **kwargs)
```

- Create a new database session and return a new connection object.
- The connection parameters can be specified as a libpq connection string using the dsn parameter:

```
conn = psycopg2.connect("dbname=test user=postgres password=secret")
```

- or using a set of keyword arguments:

```
conn = psycopg2.connect(dbname="test", user="postgres", password="secret")
```

connect()

- The basic connection parameters are:
 - dbname – the database name (database is a deprecated alias)
 - user – user name used to authenticate
 - password – password used to authenticate
 - host – database host address (defaults to UNIX socket if not provided)
 - port – connection port number (defaults to 5432 if not provided)

The connection class

- Handles the connection to a PostgreSQL database instance. It encapsulates a database session.
- Connections are created using the factory function `connect()`.
- Connections are thread safe and can be shared among many threads. See Thread and process safety for details.

connection.cursor() method

`cursor(name=None, cursor_factory=None, scrollable=None, withhold=False)`

- Return a new cursor object using the connection.
- If name is specified, the returned cursor will be a server side cursor (also known as named cursor). Otherwise it will be a regular client side cursor.

`connection.commit(), connection.rollback(), connection.close()`

commit()

- Commit any pending transaction to the database.
- By default, Psycopg opens a transaction before executing the first command: if `commit()` is not called, the effect of any data manipulation will be lost.
- The connection can be also set in “autocommit” mode: no transaction is automatically open, commands have immediate effect.

rollback()

- Roll back to the start of any pending transaction. Closing a connection without committing the changes first will cause an implicit rollback to be performed.

close()

- Close the connection now. The connection will be unusable from this point forward

The cursor class

- Allows Python code to execute PostgreSQL command in a database session. Cursors are created by the `connection.cursor()` method: they are bound to the connection for the entire lifetime and all the commands are executed in the context of the database session wrapped by the connection.
- Cursors created from the same connection are not isolated, i.e., any changes done to the database by a cursor are immediately visible by the other cursors. Cursors created from different connections can or can not be isolated, depending on the connections' isolation level.

close() method

- Close the cursor now. The cursor will be unusable from this point forward; an `InterfaceError` will be raised if any operation is attempted with the cursor.

Commands execution methods

execute(query, vars=None)

- Execute a database operation (query or command).
- Parameters may be provided as sequence or mapping and will be bound to variables in the operation. Variables are specified either with positional (%s) or named %(name)s placeholders.
- The method returns None. If a query was executed, the returned values can be retrieved using fetch*() methods.

executemany(query, vars_list)

- Execute a database operation (query or command) against all parameter tuples or mappings found in the sequence vars_list.
- The function is mostly useful for commands that update the database: any result set returned by the query is discarded.
- Parameters are bounded to the query using the same rules described in the execute() method.

```
>>> nums = ((1,), (5,), (10,))
>>> cur.executemany("INSERT INTO test (num) VALUES (%s)", nums)

>>> tuples = ((123, "foo"), (42, "bar"), (23, "baz"))
>>> cur.executemany("INSERT INTO test (num, data) VALUES (%s, %s)", tuples)
```

Results retrieval methods

- The following methods are used to read data from the database after an execute() call:

fetchone()

- Fetch the next row of a query result set, returning a single tuple, or None when no more data is available:

```
>>> cur.execute("SELECT * FROM test WHERE id = %s", (3,))
>>> cur.fetchone()
(3, 42, 'bar')
```

- A programming error is raised if the previous call to execute*() did not produce any result set or no call was issued yet.

Results retrieval methods

fetchall()

- Fetch all (remaining) rows of a query result, returning them as a list of tuples. An empty list is returned if there is no more record to fetch.

```
>>> cur.execute("SELECT * FROM test;")
>>> cur.fetchall()
[(1, 100, "abc'def"), (2, None, 'dada'), (3, 42, 'bar')]
```

- A programming error is raised if the previous call to `execute*()` did not produce any result set or no call was issued yet.

Adaptation of Python values to SQL types

- Many standard Python types are adapted into SQL and returned as Python objects when a query is executed.
- The following table shows the default mapping between Python and PostgreSQL types:

Python	PostgreSQL
None	NULL
bool	bool
float	real double
int long	smallint integer bigint
Decimal	numeric
str unicode	varchar text
buffer memoryview bytearray bytes Buffer protocol	bytea

date	date
time	time timetz
datetime	timestamp timestampz
timedelta	interval
list	ARRAY
tuple namedtuple	Composite types IN syntax
dict	hstore
Psycopg's Range	range
Anything™	json
UUID	uuid
ipaddress objects	inet cidr

Further study

<https://www.psycopg.org/docs/index.html>

Visualization of spatial data with Matplotlib and Python

- Matplotlib is a powerful and very popular data visualization library in Python.
- Documentation and examples available at:
<https://matplotlib.org/stable/index.html>




Creating a scatter plot to show the taxi stands of Porto

```
import matplotlib.pyplot as plt
import psycopg2

scale=1/30000
conn = psycopg2.connect("dbname=michelferreira user=michelferreira")
cursor_psycopg2 = conn.cursor()
sql = "select st_astext(proj_location) from taxi_stands"
cursor_psycopg2.execute(sql)
results = cursor_psycopg2.fetchall()
xs = []
ys = []
for row in results:
    point_string = row[0]
    point_string = point_string[6:-1]
    (x,y) = point_string.split()
    xs.append(float(x))
    ys.append(float(y))
width_in_inches = ((max(xs)-min(xs))/0.0254)*1.1
height_in_inches = ((max(ys)-min(ys))/0.0254)*1.1
fig = plt.figure(figsize=(width_in_inches*scale,height_in_inches*scale))
plt.scatter(xs,ys,s=5)
plt.show()
```


Creating a scatter plot to show the taxi stands of Porto



Creating a line plot with the boundary of the convex hull of taxi stands


```
sql = "select st_astext(st_convexhull(st_collect(proj_location))) from taxi_stands"
cursor_psql.execute(sql)
results = cursor_psql.fetchall()

xs = []
ys = []

for row in results:
    print(row[0])
    points_string = row[0]
    points_string = points_string[9:-2]
    points = points_string.split(',')
    for point in points:
        (x,y) = point.split()
        xs.append(float(x))
        ys.append(float(y))
        print(x,y)

plt.plot(xs,ys, color='red')
plt.show()
```

Creating a line plot with the boundary of the convex hull of taxi stands



Visualizing the parishes of the district of Porto

